

Partial Information Decomposition via Deficiency for Multivariate Gaussians

Gabriel Schamberg*
gabes@mit.edu

Picower Institute for Learning and Memory
Massachusetts Institute of Technology

Praveen Venkatesh*
vpraveen@cmu.edu

Dept. of Electrical and Computer Engineering
Carnegie Mellon University

Abstract—We consider the problem of decomposing the joint information content of three jointly Gaussian random vectors using the partial information decomposition (PID). Specifically, we extend previous works characterizing the PID of a scalar Gaussian message by providing a necessary and sufficient condition for the existence of unique information when the message is a Gaussian random vector. This extension follows naturally from a connection between notions of stochastic degradedness of broadcast channels and Blackwell sufficiency of statistical experiments, which we discuss formally. The condition we identify indicates that the closed form decomposition for the scalar case rarely extends to the vector case. We also provide a convex optimization approach for approximating a PID in the vector case, analyze its properties, and evaluate it empirically on randomly generated Gaussian systems.

I. INTRODUCTION

Partial information decompositions (PIDs) provide a framework for characterizing the joint information content of three or more random variables. The three-variable case is usually discussed in terms of how the information content in a “message” variable, M , is decomposed between *two* others, X and Y , and is hence commonly referred to as the *bivariate* decomposition. More precisely, a bivariate PID decomposes the mutual information between M and the pair (X, Y) into four components: the information about M that is (i) unique to X ; (ii) unique to Y ; (iii) redundant in both X and Y ; and (iv) synergistic, i.e., cannot be obtained from X or Y *individually*, but is present in their combination (X, Y) . This work focuses on the bivariate PID for jointly Gaussian random vectors M , X and Y .

Recently, several competing definitions have emerged for the general bivariate case [1–6] (see [7] for a review). Some PID definitions, which we believe are operationally better motivated than others, have an important feature in common: they satisfy (what we call) the “Blackwell criterion” [8]. Informally, this criterion states that X has unique information about M if and only if, on average, one can make better decisions about M (w.r.t. any loss function) by observing X rather than Y . For a few of the definitions that satisfy the Blackwell criterion, the Gaussian case was previously examined by Barrett [9]. Barrett showed that when M is restricted to be scalar (while X and Y may be vectors), the Blackwell criterion is *always* satisfied, i.e., only one of X or Y can have unique information about M . This result also

provides a straightforward way to *compute* the PID in such cases, since the redundant information is easily shown to be the minimum of $I(M; X)$ and $I(M; Y)$. However, the case of *fully* multivariate Gaussians¹ has remained unresolved in the PID literature: (1) *When does the Blackwell criterion hold?* and (2) *How do we compute the PID when it does not?*

Our work extends Barrett’s result to the case of fully multivariate Gaussian random variables, by characterizing when the Blackwell criterion holds and providing a heuristic optimization framework for estimating the Gaussian PID when it does not. Our first result emerges from the observation that the Blackwell criterion is equivalent to the idea of stochastic degradedness of broadcast channels (an observation also made by Raginsky [10]). We then draw on the results of Gerdes et al. [11], who (building on the previous work of Shang and Poor [12]) provide a necessary and sufficient condition for the stochastic degradedness of linear additive Gaussian-noise broadcast channels. By extension, this is also a necessary and sufficient condition for the Blackwell criterion for fully multivariate Gaussians. It should be noted that this result was also shown much earlier in the statistical decision theory literature using a different proof technique [13, Thm. 8.2.13].

The central contribution of our extension of Barrett’s result lies, therefore, in tying together disparate works from different literatures and arriving at an implication for the Gaussian PID, which was hitherto unknown. Indeed, our work shows that Barrett’s result is simply a partial extension of the well-known fact that scalar Gaussian broadcast channels are always stochastically degraded [14, Example 15.6.6].

The characterization of the Blackwell criterion also inspires a convex optimization framework for *computing* a deficiency-based PID definition [4] for fully multivariate Gaussians. The notion of *deficiency*, first introduced by Le Cam [15] (see [13] for a thorough exposition), quantifies the lack of Blackwell *sufficiency*, and so is a natural PID measure. We show that the deficiency computed using our optimization framework satisfies some basic desirable properties, but that several open questions remain; we also examine a few of these empirically.

In Section II we provide background on the basics of the PID and its relationship to the Blackwell criterion and broadcast channels. In Section III we state the equivalence of

*Equal contribution; author order was determined by a coin flip. GS was supported by the Picower Postdoctoral Fellowship.

¹We use the term “fully multivariate Gaussians” here to refer to instances of a *bivariate* PID where M may also be a vector (in contrast to Barrett’s work). This is not to be confused with *multivariate PIDs*, where information about M is decomposed among more than two variables.

Blackwell sufficiency and stochastic degradedness and present our extension of Barrett's result characterizing the existence of unique information in fully multivariate Gaussians. In Section IV we provide a convex optimization framework to approximately compute the deficiency-based PID for multivariate Gaussians. We present empirical results in Section V and conclude with a discussion of open questions in Section VI.

II. BACKGROUND

A. Partial Information Decomposition

Let M , X and Y be three random variables with sample spaces \mathcal{M} , \mathcal{X} and \mathcal{Y} respectively, and joint density P_{MXY} . Williams and Beer [1] proposed the following general decomposition for the bivariate case:

$$I(M; (X, Y)) = UI(M : X \setminus Y) + UI(M : Y \setminus X) + RI(M : X; Y) + SI(M : X; Y), \quad (1)$$

where $UI(M : X \setminus Y)$ is the information about M uniquely present in X and not in Y , while $RI(M : X; Y)$ and $SI(M : X; Y)$ are respectively the redundant and synergistic information about M contained between X and Y . For ease of notation, we will interchangeably represent these PID quantities as UI_X , UI_Y , RI , and SI respectively. Such a decomposition should also satisfy

$$I(M; X) = UI(M : X \setminus Y) + RI(M : X; Y), \quad (2)$$

$$I(M; Y) = UI(M : Y \setminus X) + RI(M : X; Y). \quad (3)$$

Based on these three equations, defining any one of the four quantities in the RHS of (1) suffices to determine the other three. For instance, Williams and Beer [1] proposed the following definition for redundancy.

Definition 1 (MMI-PID [1]). *Williams and Beer proposed that the redundant information about M present in both X and Y be given by the Minimum of their respective Mutual Informations with M (hence ‘‘MMI’’):*

$$RI_{\text{MMI}}(M : X; Y) \triangleq \min\{I(M; X), I(M; Y)\}. \quad (4)$$

The above definition, with equations (1)–(3), fully determines the MMI-PID, i.e., $UI_{\text{MMI}}(M : X \setminus Y)$, $UI_{\text{MMI}}(M : Y \setminus X)$, and $SI_{\text{MMI}}(M : X; Y)$ are now well-defined.

B. The Blackwell Criterion

Before stating the two other PID definitions we refer to in this paper, we introduce the notion of Blackwell sufficiency, which forms the operational basis motivating these definitions. The original definition of Blackwell sufficiency [8] was based on statistical decision theory, addressing whether or not one could attain the *same risk* by making decisions based on one experiment, that one could with another. However, for simplicity, we present an equivalent [8] notion of Blackwell sufficiency, that is more amenable to our setup here.

Definition 2 (Blackwell sufficiency). *We say that a channel $P_{X|M}$ is Blackwell sufficient for another channel $P_{Y|M}$ if $\exists Q_{Y|X}$ such that*

$$Q_{Y|X} \circ P_{X|M} = P_{Y|M}, \quad (5)$$

where \circ represents channel composition, i.e. $\forall m \in \mathcal{M}, y \in \mathcal{Y}$,

$$(Q_{Y|X} \circ P_{X|M})(y|m) \triangleq \int Q_{Y|X}(y|x) P_{X|M}(x|m) dx. \quad (6)$$

We denote this by $P_{X|M} \succcurlyeq_M P_{Y|M}$, or more succinctly as $X \succcurlyeq_M Y$.

Intuitively, $X \succcurlyeq_M Y$ means that we can generate a new variable Y' from X (using the stochastic transformation $Q_{Y|X}$) so that the effective channel from M to Y' is equivalent to the original channel from M to Y . The relationship between Blackwell sufficiency and the remaining PID definitions can be summarized by the following property [2]:

Property 1 (The Blackwell criterion). *For a bivariate PID of $I(M; (X, Y))$, Y has zero unique information about M with respect to X if and only if $X \succcurlyeq_M Y$.*

We now state two PID definitions that have been shown to satisfy the Blackwell criterion [2, 13] but quantify departures from Blackwell sufficiency differently.²

Definition 3 (\sim -PID [2]). *The unique information about M present in X and not in Y is given by*

$$\widetilde{UI}(M : X \setminus Y) \triangleq \min_{Q \in \Delta_P} I_Q(M; X | Y), \quad (7)$$

where $\Delta_P \triangleq \{Q_{MXY} : Q_{MX} = P_{MX}, Q_{MY} = P_{MY}\}$ and I_Q is the conditional mutual information over the joint distribution Q_{MXY} .

Definition 4 (δ -PID [4]). *The (weighted output) deficiency³ of X with respect to Y about M is defined as⁴*

$$\delta(M : Y \setminus X) \triangleq \inf_{Q_{Y|X} \in \mathcal{C}(\mathcal{Y}|\mathcal{X})} \mathbb{E}_{P_M}[D(P_{Y|M} \| Q_{Y|X} \circ P_{X|M})], \quad (8)$$

where $\mathcal{C}(\mathcal{Y}|\mathcal{X})$ is the set of all channels from \mathcal{X} to \mathcal{Y} . The deficiency-based redundant information about M present in X and Y is given by

$$RI_\delta(M : X; Y) \triangleq \min\{I(M; X) - \delta(M : X \setminus Y), I(M; Y) - \delta(M : Y \setminus X)\}. \quad (9)$$

As with the MMI-PID, equations (1)–(3) fully determine the remaining components of the \sim -PID and δ -PID. We are now in a position to formally state Barrett's result for the PID of Gaussian random variables.

Theorem 1 (Barrett [9]). *If M , X and Y are all jointly Gaussian, and M is scalar, then we always have that either $X \succcurlyeq_M Y$ or $Y \succcurlyeq_M X$. Furthermore, the \sim -PID (as well as the PID of Harder et al. [3]) is equivalent to the MMI-PID.*

We omit a proof of this theorem, since it is subsumed by the extension derived in Section III.

²So far, PIDs have mostly been defined for *discrete* random variables. Barrett [9] did not formally extend the PID definitions of other authors to continuous variables even though he analyzed Gaussians. In the definitions given here, we consider extensions to the continuous case, but relegate a detailed discussion of measure-theoretic issues to future work. For our purposes, we assume that all joint and conditional probability measures, as well as information measures thereof, are well defined.

³There are many ways to define deficiency; Raginsky [10] provides a number of these that consider the worst-case over M . We prefer an expectation over M , since M is a random variable in our setup.

⁴The reason for this notation is that the deficiency of X w.r.t. Y translates to the unique information present in Y and not in X .

C. Broadcast channels

To extend Barrett's result to the case of fully multivariate Gaussians, we will leverage the connection between Blackwell sufficiency and the concept of stochastic degradedness from the literature on broadcast channels. Following Gerdes et al. [11], we define stochastic degradedness as follows:

Definition 5 (Stochastic degradedness). *We say that a channel $P_{Y|M}$ is stochastically degraded with respect to another channel $P_{X|M}$ if there exists a random variable X' with sample space \mathcal{X} such that $P_{X'|M} = P_{X|M}$ and $M-X'-Y$ is a Markov chain.*

Raginsky [10] points out that stochastic degradedness is identical to Blackwell sufficiency, but does not offer a formal proof. We state and prove this in the following section.

III. EXISTENCE OF UNIQUE INFORMATION IN MULTIVARIATE GAUSSIANS

In this section, we show that the concepts of Blackwell sufficiency and stochastic degradedness are identical. We also characterize a necessary and sufficient condition under which stochastic degradedness (and hence the Blackwell criterion) holds for fully multivariate Gaussians. Let the sample spaces of M , X and Y respectively be $\mathcal{M} = \mathbb{R}^{d_M}$, $\mathcal{X} = \mathbb{R}^{d_X}$ and $\mathcal{Y} = \mathbb{R}^{d_Y}$. We parameterize the joint distribution P_{MXY} as:

$$M \sim P_M = \mathcal{N}(0, \Sigma_M) \quad (10)$$

$$X|M \sim P_{X|M} = \mathcal{N}(H_X M, \Sigma_X) \quad (11)$$

$$Y|M \sim P_{Y|M} = \mathcal{N}(H_Y M, \Sigma_Y) \quad (12)$$

Here, we do not explicitly define the covariance between X and Y , however we also *do not* assume that $X \perp\!\!\!\perp Y|M$. Since \widetilde{UI} and RI_δ depend only on P_M , $P_{X|M}$ and $P_{Y|M}$, our proofs will not require the use of $\text{Cov}(X, Y)$.

Remark 1. *Without loss of generality, we assume that the noise covariance matrices Σ_X and Σ_Y are full rank. This can always be done, since removing linearly dependent elements within X and Y respectively has no effect on information-theoretic relationships.*

First, it is worth noting that Barrett's result does *not* directly extend to the case of jointly Gaussian variables when M is a vector. Barrett appears to leave this as an open question [9, end of Sec. IV-B], but it can be answered with a simple counterexample:

Counterexample. Consider the case where $M = [M_1, M_2]$, $X = M_1 + \epsilon N_X$ and $Y = M_2 + \epsilon N_Y$, with $M_1, M_2, N_X, N_Y \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$; $\epsilon \approx 0$. Intuitively, X and Y each have equal amounts of unique information about M , and one can show that this holds for the \sim - and δ -PIDs. However, under the MMI-PID, X and Y have only redundant and synergistic information and no unique information. \square

Before moving to our main result, we formally connect Blackwell sufficiency and stochastic degradedness:

Lemma 2 (Equivalence of Blackwell sufficiency and stochastic degradedness). *$P_{X|M}$ is Blackwell sufficient for $P_{Y|M}$ if and only if $P_{Y|M}$ is stochastically degraded with respect to $P_{X|M}$.*

The proof is quite trivial and primarily involves juggling notation, but is provided in Appendix A for completeness. In fact, Cover and Thomas [14, Sec. 15.6.2] even *define* stochastic degradedness similar to how we have defined Blackwell sufficiency. We are now in a position to state the main result of this paper.

Theorem 3. *For jointly Gaussian random vectors M , X and Y , $X \succsim_M Y$ if and only if $H_X^T \Sigma_X^{-1} H_X \succsim H_Y^T \Sigma_Y^{-1} H_Y$.*

The proof of this result follows directly from the result of Gerdes et al. [11], along with Lemma 2. In an effort to present a self-contained record of proofs, we provide it in Appendix B. The main result can be interpreted as follows: for jointly Gaussian vectors M , X and Y , and a PID satisfying the Blackwell criterion, X and Y *both* have unique information about M *unless* the condition in Theorem 3 holds (in one direction or the other). This is significant, as the condition is fairly restrictive: when $d_X, d_Y < d_M$, the condition will hold w.p. 0 assuming a continuous distribution over H_X and H_Y [11, Prop. 2]. Given that the MMI-PID assigns unique information to only one of X or Y , it is clear that an alternative approach to computing PIDs is needed when $d_M > 1$.

Remark 2. *Since Σ_X and Σ_Y are invertible (Remark 1), we can equivalently write the condition in Theorem 3 using a whitening transform:*

$$\tilde{H}_X^T \tilde{H}_X \succsim \tilde{H}_Y^T \tilde{H}_Y; \quad \tilde{H}_X \triangleq \Sigma_X^{-\frac{1}{2}} H_X, \quad \tilde{H}_Y \triangleq \Sigma_Y^{-\frac{1}{2}} H_Y \quad (13)$$

Without loss of generality, we assume henceforth that H_X and H_Y have been whitened, and hence $\Sigma_X = \Sigma_Y = I$.

IV. APPROXIMATING THE δ -PID

Having established a complete characterization of when unique information exists in the multivariate Gaussian setting, we now seek to compute the δ -PID. This requires computing the weighted output deficiency given by (8), which is challenging because it requires an optimization over all continuous conditional distributions $Q_{Y|X}$. Instead, we consider the restricted problem where $Q_{Y|X}$ lies in the set of linear additive Gaussian noise channels $\mathcal{C}_G(\mathcal{Y}|\mathcal{X}) \subset \mathcal{C}(\mathcal{Y}|\mathcal{X})$. Then, $Q_{Y|X}$ can be parameterized in terms of its channel gain and noise covariance matrices, $T \in \mathbb{R}^{d_Y \times d_X}$ and $\Sigma_T \in \mathbb{R}^{d_Y \times d_Y}$, $\Sigma_T \succ 0$, so that the *Gaussian deficiency* can be defined as:

$$\begin{aligned} \delta_G(M : Y \setminus X) &\triangleq \inf_{Q_{Y|X} \in \mathcal{C}_G(\mathcal{Y}|\mathcal{X})} \mathbb{E}_{P_M} [D(P_{Y|M} \| Q_{Y|X} \circ P_{X|M})] \quad (14) \end{aligned}$$

$$\begin{aligned} &= \inf_{T, \Sigma_T \succ 0} \frac{1}{2} \left[\mathbb{E}_{P_M} \left[\left\| (TH_X - H_Y)M \right\|_{\Sigma_T + T\Sigma_X T^T}^2 \right] \right. \\ &\quad \left. + \text{Tr}((\Sigma_T + T\Sigma_X T^T)^{-1} \Sigma_Y) \right. \\ &\quad \left. + \log \frac{|\Sigma_T + T\Sigma_X T^T|}{|\Sigma_Y|} - d_Y \right] \quad (15) \end{aligned}$$

$$\begin{aligned} &= \inf_{T, \Sigma_T \succ 0} \frac{1}{2} \left[\mathbb{E}_{P_M} \left[\left\| (TH_X - H_Y)M \right\|_{\Sigma_T + TT^T}^2 \right] \right. \\ &\quad \left. + \text{Tr}((\Sigma_T + TT^T)^{-1}) \right. \\ &\quad \left. + \log |\Sigma_T + TT^T| - d_Y \right], \quad (16) \end{aligned}$$

where $\|a\|_B^2 \triangleq aB^{-1}a^T$ is the squared Mahalanobis distance, and (16) follows from Remark 2. Unfortunately, this problem may not be convex since $\text{Tr}((\Sigma_T + TT^T)^{-1})$ is the composition of a convex function ($\text{Tr}((\cdot)^{-1})$) with a non-monotonic function of T . Therefore, we propose a convex relaxation of the problem in (14) to find an approximate minimizer $\hat{Q}_{Y|X} = \mathcal{N}(\hat{T}X, \hat{\Sigma}_T)$. To motivate the proposed relaxation, we note the deficiency is bounded as $0 \leq \delta_G(M : Y \setminus X) \leq I(M; Y)$ and target our approach to recover the true deficiency when it takes these extremal values. Letting $Q_{Y|X}^* = \mathcal{N}(T^*X, \Sigma_T^*)$ be the minimizer of (14), we note that

$$\begin{aligned} T^*H_X &= H_Y, \quad \Sigma_T^* = I - T^*T^{*T} \Rightarrow \delta_G(M : Y \setminus X) = 0; \\ T^* &= \mathbf{0}, \quad \Sigma_T^* = I + H_Y \Sigma_M H_Y^T \Rightarrow \delta_G(M : Y \setminus X) = I(M; Y). \end{aligned}$$

Intuitively, as the ability to construct a channel $Q_{Y|X} \circ P_{X|M}$ that replicates $P_{Y|M}$ diminishes (i.e. the deficiency approaches the mutual information), we need to begin incorporating the noise associated with the marginal distribution P_Y in $\hat{\Sigma}_T$. Thus, given an estimated \hat{T} , we construct $\hat{\Sigma}_T$ as:

$$\hat{\Sigma}_T = I + H_Y \Sigma_M H_Y^T - \hat{T}(I + H_X \Sigma_M H_X^T) \hat{T}^T \quad (17)$$

The remaining problem is identifying a reasonable gain matrix \hat{T} for $\hat{Q}_{Y|X}$. Since we cannot directly minimize the first term in (16), we formulate a convex relaxation to approximate T^* :

$$\begin{aligned} \hat{T} &= \underset{T}{\text{argmin}} \quad \mathbb{E}_{P_M} \left[\|(TH_X - H_Y)M\|_{I+H_Y \Sigma_M H_Y^T}^2 \right] \\ \text{s.t.} \quad \Sigma_T &\succcurlyeq 0 \end{aligned} \quad (18)$$

with Σ_T as defined in (17) with T in place of \hat{T} . The choice to minimize the Mahalanobis distance w.r.t. $I + H_Y \Sigma_M H_Y^T$ is motivated in the proof of Proposition 4 and is partially justified by the empirical results in Section V, with further justifications and/or demonstrations of optimality left for future work.

Having defined an approach for obtaining $\hat{Q}_{Y|X}$, the resulting approximate deficiency is given by:

$$\begin{aligned} \hat{\delta}_G(M : Y \setminus X) &\triangleq \frac{1}{2} \left[\mathbb{E}_{P_M} \left[\left\| (\hat{T}H_X - H_Y)M \right\|_{\hat{\Sigma}_T + \hat{T}\hat{T}^T}^2 \right] \right. \\ &\quad \left. + \text{Tr}((\hat{\Sigma}_T + \hat{T}\hat{T}^T)^{-1}) \right. \\ &\quad \left. + \log |\hat{\Sigma}_T + \hat{T}\hat{T}^T| - d_Y \right] \end{aligned} \quad (19)$$

It can be shown that $\hat{\Sigma}_T + \hat{T}\hat{T}^T$ is guaranteed to be invertible when using the relaxation of (17) and (18). We state this formally and prove it in Appendix C. Because $\hat{Q}_{Y|X} \in \mathcal{C}_G(Y|X) \subset \mathcal{C}(Y|X)$, any method for obtaining \hat{T} and $\hat{\Sigma}_T$ will yield:

$$\hat{\delta}_G(M : Y \setminus X) \geq \delta_G(M : Y \setminus X) \geq \delta(M : Y \setminus X) \quad (20)$$

Now we provide conditions for when the estimate in (19) coincides with the true deficiency:

Proposition 4. *For jointly Gaussian random vectors M , X and Y , the approximate deficiency defined by (17)–(19) satisfies:*

$$\delta(M : Y \setminus X) = 0 \Leftrightarrow \hat{\delta}_G(M : Y \setminus X) = 0 \quad (21)$$

$$\delta(M : Y \setminus X) = I(M; Y) \Rightarrow \hat{\delta}_G(M : Y \setminus X) = I(M; Y) \quad (22)$$

The proof is given in Appendix D. Plugging in $\hat{\delta}_G$ into (9), we obtain an approximation of RI_δ , and hence an approximation of the δ -PID, which we call the $\hat{\delta}_G$ -PID:

$$\begin{aligned} \widehat{RI}(M : X; Y) &\triangleq \min \{ I(M; X) - \hat{\delta}_G(M : X \setminus Y), \\ &\quad I(M; Y) - \hat{\delta}_G(M : Y \setminus X) \} \end{aligned} \quad (23)$$

For brevity, we use \widehat{RI} etc. when referring to the constituent atoms of the $\hat{\delta}_G$ -PID. As with other PIDs, \widehat{RI} paired with equations (1)–(3) fully determines the remaining atoms of the $\hat{\delta}_G$ -PID: \widehat{UI}_X , \widehat{UI}_Y and \widehat{SI} . Conveniently, the inequalities in (20) bound the approximated $\hat{\delta}_G$ -PID by the true δ -PID:

Proposition 5. *For jointly Gaussian random vectors M , X and Y , the approximate δ -PID satisfies:*

$$\widehat{RI}(M : X; Y) \leq RI_\delta(M : X; Y) \quad (24)$$

$$\widehat{UI}(M : X \setminus Y) \geq UI_\delta(M : X \setminus Y) \quad (25)$$

$$\widehat{UI}(M : Y \setminus X) \geq UI_\delta(M : Y \setminus X) \quad (26)$$

$$\widehat{SI}(M : X; Y) \leq SI_\delta(M : X; Y) \quad (27)$$

The proof follows directly by applying the inequality in (20) to compare (9) and (23), and from the basic PID equations. It is worth noting that this proposition holds for *any* \hat{T} and $\hat{\Sigma}_T \succcurlyeq 0$ (not just the minimizer of (18)), and is thus of limited utility as a standalone result. For example, we could simply define $\hat{T} = \mathbf{0}$ and $\hat{\Sigma}_T = I + H_Y \Sigma_M H_Y^T$, and the proposition would recover trivial bounds implied by the basic PID equations ($RI_\delta \geq 0$, $UI_{\delta,X} \leq I(M; X)$, $UI_{\delta,Y} \leq I(M; Y)$, and $SI_\delta \geq I(M; (X, Y)) - I(M; X) - I(M; Y)$). However, we believe Proposition 5 is worth stating for two reasons. First, it highlights that our proposed approach is in some sense complementary to the MMI-PID, as the MMI-PID recovers trivial bounds in the opposite direction, i.e. upper bounds on redundant and synergistic information and lower bounds on unique information. Second, if the approximate δ -PID is a valid PID (i.e., if all atoms are non-negative and satisfy equations (1)–(3)) then each atom of the approximate δ -PID is bounded *between* the true δ -PID and the \sim -PID [2]. This is because the \sim -PID upper bounds the unique information and lower bounds the redundant and synergistic information [2, Lemma 3] of every other PID definition which satisfies a basic property⁵—one that is satisfied by the $\hat{\delta}_G$ -PID.⁶ Unfortunately, proving that the $\hat{\delta}_G$ -PID is a valid PID requires proving that $\widehat{SI} \geq 0$. We have not been able to demonstrate formally, but in the following section we test it empirically.

V. EMPIRICAL RESULTS

Having established a framework for approximating the δ -PID, we seek to address three questions using simulations:

⁵This property, known as Assumption (*) in [2], requires that the redundant and unique informations depend only on the marginal distribution P_M , and the channels $P_{X|M}$ and $P_{Y|M}$. This is easily seen to be satisfied by the δ -PID, as well as the $\hat{\delta}_G$ -PID.

⁶This line of reasoning assumes that Lemma 3 from [2] holds for the \sim -PID defined for *continuous* random variables. This assumption was implicit in the demonstration of Barrett's central result in [9], but it has not been explicitly verified to our knowledge.

(Q1) Does $\hat{\delta}_G$ yield a valid PID for jointly Gaussian systems? (Q2) Is the $\hat{\delta}_G$ -PID consistent with the results of Barrett [9] for the case of univariate M ? (Q3) How does information decompose in multivariate Gaussian systems?

To address these questions, we sample a random covariance matrix $\Sigma \in \mathbb{R}^{d \times d}$ (where $d = d_M + d_X + d_Y$) from a standard Wishart distribution, which fully characterizes the joint distribution P_{MXY} . To see how the PID varies with different dimensionalities d_M , d_X , and d_Y , we sample them using one of four sampling schemes:

- (S1) $d_M \sim \text{Unif}\{1 \dots 10\}$ and $d_X = d_Y = d_M$
- (S2) $d_M \sim \text{Unif}\{1 \dots 9\}$ and $d_X, d_Y \stackrel{\text{iid}}{\sim} \text{Unif}\{d_M + 1 \dots 10\}$
- (S3) $d_M \sim \text{Unif}\{2 \dots 10\}$ and $d_X, d_Y \stackrel{\text{iid}}{\sim} \text{Unif}\{1 \dots d_M - 1\}$
- (S4) $d_M \sim \text{Unif}\{2 \dots 9\}$, $d_X \sim \text{Unif}\{1 \dots d_M - 1\}$, and $d_Y \sim \text{Unif}\{d_M + 1 \dots 10\}$

Without loss of generality, we assume $d_X \leq d_Y$ (i.e. we switch their values if $d_Y < d_X$). For each of these sampling schemes, we consider 20,000 samples of d_M , d_X , d_Y and Σ , yielding 80,000 total systems and PID approximations. Further details on the implementation and experimental setup are provided in Appendices E and F and all of the associated code is provided at <https://github.com/gabeschamberg/mvar-gauss-pid>.

The first noteworthy result is that all of the approximated PID quantities were non-negative for every sampled system. This suggests an affirmative answer to (Q1), perhaps with the exception of a zero-measure set of systems. Our results also suggest an affirmative answer to (Q2), as every system for which $d_M = 1$ ($N = 4256$) resulted in our approximation assigning unique information to only one of X or Y .

To address (Q3), we visualize how the distribution of unique, redundant, and synergistic information changes for the different relationships between d_M , d_X , and d_Y in Figure 1. Given that the scale of the PID components varies with $I(M; (X, Y))$, we consider normalized PID quantities \overline{UI}_X , \overline{UI}_Y , \overline{RI} , and \overline{SI} obtained by dividing the corresponding approximate PID quantities by $I(M; (X, Y))$. Thus, these normalized values are all non-negative and satisfy $\overline{UI}_X + \overline{UI}_Y + \overline{RI} + \overline{SI} = 1$. In Figure 1 we represent each system by its location on a 4-simplex characterizing the proportion of $I(M; (X, Y))$ that is accounted for by each PID component.

Figure 1 provides numerous insights into the relative prevalence of the PID components in multivariate Gaussian systems. First, synergy has a strong presence in all four panels. Synergy appears to be most prevalent when $d_M \leq d_X, d_Y$ (S1 and S2). This result is not entirely surprising—the prevalence of synergy in systems where $d_M = 1$ was acknowledged by Barrett [9]. Correspondingly, (S2), which gives the closest multivariate analogue to the univariate case in the sense that $d_M = 1 \Rightarrow d_M \leq d_X, d_Y$, shows the largest synergy. Redundancy is never particularly prevalent—this may be a consequence of how we have drawn a random covariance matrix, which reduces the likelihood that X and Y capture similar “directions” of M . In other words, since each dimension of X and Y represents a linear combination of the dimensions of M plus i.i.d. Gaussian noise (after whitening),

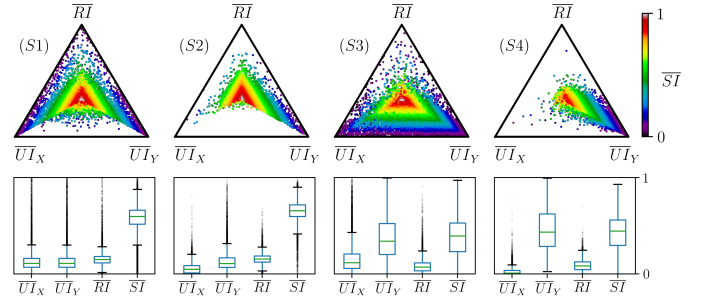


Fig. 1: Distribution of normalized unique, redundant, and synergistic information for Gaussian systems sampled from the standard Wishart distribution: (S1) $d_M = d_X = d_Y$; (S2) $d_M < d_X \leq d_Y$; (S3) $d_X \leq d_Y < d_M$; and (S4) $d_X < d_M < d_Y$. Top row: scatter plots of the computed \overline{UI}_X , \overline{UI}_Y , \overline{RI} and \overline{SI} on the 4-simplex (3D views in Appendix G). Bottom row: box plots showing the relative prevalence of each partial information atom.

the redundant information loosely captures the extent to which the measurement vectors associated with X are aligned with those of Y . Under this interpretation, we would expect there to be less alignment when there are fewer vectors (d_X, d_Y are small) of larger dimensionality (d_M is large). That may be why (S3) and (S4) have the least redundancy, since X has very few dimensions to begin with in both these cases. Finally, because we have set $d_Y \geq d_X$, we see that Y tends to have more unique information than X . The only exception is that of (S1), where $d_X = d_Y$. (S3) and (S4) have large amounts of unique information in Y , since d_M is large in both these cases, implying that there are more dimensions of M that X and Y might uniquely capture. This effect is particularly pronounced in (S4), as Y has enough dimensions to capture M , while X rarely captures any part of M uniquely.

VI. DISCUSSION AND OPEN QUESTIONS

While we believe that our proposed approximation provides the best existing method for computing bounds on the PID in multivariate Gaussian systems, there are shortcomings with the proposed approach. Most importantly, we have not been able to prove that our approximate PID is always non-negative. This is not a huge issue in practice, as negative estimates of PID quantities can be discarded and replaced by the bounds implied by the basic PID equations, but proving that the approximate PID is a valid PID could have implications for bounding it between the \sim -PID and δ -PID. A second shortcoming is that the $\hat{\delta}_G$ -PID is not guaranteed to satisfy the Blackwell criterion as a result of not having an upper bound on the approximate deficiency. The simulation results are favorable for both of these issues, but formal proofs are still desirable.

A number of questions regarding the computation of deficiency remain, including whether it can be computed exactly and conditions under which (8) and (14) are equivalent. Estimating the Gaussian PID for practical applications requires estimating the full covariance matrix Σ , so it will be important to develop statistical tests for characterizing confidence in PID approximations. Finally, another future direction involves extending the relationship between stochastic degradedness and unique information from Gaussian channels to general continuous channels [16].

REFERENCES

- [1] P. L. Williams and R. D. Beer, "Nonnegative decomposition of multivariate information," *arXiv preprint arXiv:1004.2515*, 2010.
- [2] N. Bertschinger, J. Rauh, E. Olbrich, J. Jost, and N. Ay, "Quantifying unique information," *Entropy*, vol. 16, no. 4, pp. 2161–2183, 2014.
- [3] M. Harder, C. Salge, and D. Polani, "Bivariate measure of redundant information," *Physical Review E*, vol. 87, no. 1, p. 012130, 2013.
- [4] P. K. Banerjee, E. Olbrich, J. Jost, and J. Rauh, "Unique informations and deficiencies," in *2018 56th Annual Allerton Conference on Communication, Control, and Computing (Allerton)*. IEEE, 2018, pp. 32–38.
- [5] X. Niu and C. J. Quinn, "A measure of synergy, redundancy, and unique information using information geometry," in *2019 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2019, pp. 3127–3131.
- [6] C. Finn and J. T. Lizier, "Pointwise partial information decomposition using the specificity and ambiguity lattices," *Entropy*, vol. 20, no. 4, p. 297, 2018.
- [7] J. T. Lizier, N. Bertschinger, J. Jost, and M. Wibral, "Information decomposition of target effects from multi-source interactions: perspectives on previous, current and future work," p. 307, 2018.
- [8] D. Blackwell, "Equivalent comparisons of experiments," *The Annals of Mathematical Statistics*, pp. 265–272, 1953.
- [9] A. B. Barrett, "Exploration of synergistic and redundant information sharing in static and dynamical Gaussian systems," *Physical Review E*, vol. 91, no. 5, p. 052802, 2015.
- [10] M. Raginsky, "Shannon meets Blackwell and Le Cam: Channels, codes, and statistical experiments," in *2011 IEEE International Symposium on Information Theory Proceedings*. IEEE, 2011, pp. 1220–1224.
- [11] L. Gerdes, M. Riemensberger, and W. Utschick, "On the equivalence of degraded Gaussian MIMO broadcast channels," in *WSA 2015; 19th International ITG Workshop on Smart Antennas*. VDE, 2015, pp. 1–5.
- [12] X. Shang and H. V. Poor, "Noisy-interference sum-rate capacity for vector Gaussian interference channels," *IEEE transactions on information theory*, vol. 59, no. 1, pp. 132–153, 2012.
- [13] E. Torgersen, *Comparison of statistical experiments*. Cambridge University Press, 1991, vol. 36.
- [14] T. M. Cover and J. A. Thomas, *Elements of Information Theory*. John Wiley & Sons, 2012.
- [15] L. Le Cam, "Sufficiency and approximate sufficiency," *Ann. Math. Statist.*, vol. 35, no. 4, pp. 1419–1455, 12 1964. [Online]. Available: <https://doi.org/10.1214/aoms/1177700372>
- [16] A. Makur and Y. Polyanskiy, "Comparison of channels: Criteria for domination by a symmetric channel," *IEEE Transactions on Information Theory*, vol. 64, no. 8, pp. 5704–5725, 2018.
- [17] F. Zhang, *The Schur complement and its applications*. Springer Science & Business Media, 2006, vol. 4.

APPENDIX A

PROOF OF LEMMA 2

Proof of Lemma 2. Throughout this proof, we drop the arguments of probability distributions. When we equate two distributions, we mean that they are identical at all points in their shared domain.

(\Leftarrow) Suppose $P_{Y|M}$ is stochastically degraded w.r.t. $P_{X|M}$. Then, $\exists X'$ such that

$$P_{X'|M} = P_{X|M} \text{ and } P_{YX'|M} = P_{Y|X'}P_{X'|M}. \quad (28)$$

Let $Q_{Y'|X} \triangleq P_{Y|X}$. Then, (28) implies that

$$Q_{Y'|X}P_{X|M} = P_{Y|X'}P_{X'|M}. \quad (29)$$

Therefore,

$$\int Q_{Y'|X}P_{X|M} dx = \int P_{Y|X'}P_{X'|M} dx \quad (30)$$

$$= \int P_{YX'|M} dx = P_{Y|M}. \quad (31)$$

(\Rightarrow) Suppose $P_{X|M}$ is Blackwell sufficient for $P_{Y|M}$. Then, $\exists Q_{Y'|M}$ such that

$$\int Q_{Y'|X}P_{X|M} dx = P_{Y|M}. \quad (32)$$

In other words, if we let Q represent the joint distribution between M , X and Y' , we have that

$$Q_{Y'|M} = P_{Y|M} \text{ and hence } Q_{Y'M} = P_{Y'M}. \quad (33)$$

Let X' be defined through a stochastic transformation of Y and M : $P_{X'|Y'M}(x|y, m) \triangleq Q_{X|Y'M}(x|y, m) \forall x \in \mathbf{X}, y \in \mathbf{Y}, m \in \mathbf{M}$. Then, using (33), we find

$$P_{X'Y'M} = P_{X'|Y'M}P_{Y'M} \quad (34)$$

$$= Q_{X|Y'M}Q_{Y'M} = Q_{XY'M}. \quad (35)$$

This in turn implies

$$P_{X'Y|M} = Q_{XY'M} \stackrel{(a)}{=} Q_{Y'|X}Q_{X|M} \quad (36)$$

$$\stackrel{(b)}{=} P_{Y|X'}P_{X'|M}, \quad (37)$$

which proves that $M-X'-Y$ is a Markov chain. In the above equation, (a) follows from how $Q_{Y'|X}$ is defined, while (b) follows from (34). \square

APPENDIX B

PROOF OF THEOREM 3

Before proceeding to the proof, we introduce a lemma borrowed from Shang and Poor [12, Lemma 5] that provides an equivalent characterization of the condition in Theorem 3.

Lemma 6. *The condition $H_X^T \Sigma_X^{-1} H_X \succcurlyeq H_Y^T \Sigma_Y^{-1} H_Y$ holds if and only if*

$$\exists T \text{ s.t. } H_Y = TH_X \text{ and } T\Sigma_X T^T \preccurlyeq \Sigma_Y. \quad (38)$$

Proof. Consider the whitened form of the channels:

$$\tilde{H}_X \triangleq \Sigma_X^{-\frac{1}{2}} H_X, \tilde{H}_Y \triangleq \Sigma_Y^{-\frac{1}{2}} H_Y. \quad (39)$$

Then, we need to show $\tilde{H}_X^T \tilde{H}_X \succcurlyeq \tilde{H}_Y^T \tilde{H}_Y$ if and only if

$$\exists T \text{ s.t. } \tilde{H}_Y = T\tilde{H}_X \text{ and } TT^T \preccurlyeq I. \quad (40)$$

The remainder of the proof follows from [12, Lemma 5]. \square

Proof of Theorem 3. This proof is derived in large part from the work of Gerdes et al. [11].

(\Leftarrow) Suppose that $H_X^T \Sigma_X^{-1} H_X \succcurlyeq H_Y^T \Sigma_Y^{-1} H_Y$. Then, by Lemma 6, $\exists T$ such that $H_Y = TH_X$ and $T\Sigma_X T^T \preccurlyeq \Sigma_Y$. In other words, we may write

$$Y' = TX + N, \quad (41)$$

where $N \sim \mathcal{N}(0, \Sigma_T)$, with $\Sigma_T \triangleq \Sigma_Y - T\Sigma_X T^T$ (the fact that $T\Sigma_X T^T \preccurlyeq \Sigma_Y$ ensures that Σ_T is positive semidefinite, and hence a valid covariance matrix). Therefore, $\exists Y'$ generated by a stochastic transformation $P_{Y'|X} = \mathcal{N}(TX, \Sigma_T)$, such that

$$P_{Y'|M} = \mathcal{N}(TH_X M, T\Sigma_X T^T + \Sigma_T) \quad (42)$$

$$= \mathcal{N}(H_Y M, \Sigma_Y) = P_{Y|M}. \quad (43)$$

Hence, by Definition 2, $X \succcurlyeq_M Y$.

(\Rightarrow) Next, suppose that $X \succcurlyeq_M Y$. Then, $\exists Y'$ generated by $P_{Y'|X}$, such that $P_{Y'|M} = P_{Y|M}$. Since $M-X-Y'$ is a Markov chain,

$$I(M; X) \stackrel{(a)}{\geq} I(M; Y') \stackrel{(b)}{=} I(M; Y), \quad (44)$$

where (a) follows from the Data Processing Inequality [14, Ch. 2], and (b) follows because $P_{Y'|M} = P_{Y|M}$. For a Gaussian channel $P_{X|M}$, the mutual information is given by [14, Thm. 8.4.1]

$$I(M; X) = h(X) - h(X|M) \quad (45)$$

$$= \frac{1}{2} \log |(2\pi e)(\Sigma_X + H_X \Sigma_M H_X^T)| - \frac{1}{2} \log |2\pi e \Sigma_X| \quad (46)$$

$$= \frac{1}{2} \log \left(\frac{|\Sigma_X + H_X \Sigma_M H_X^T|}{|\Sigma_X|} \right) \quad (47)$$

$$= \frac{1}{2} \log \left(\frac{|\Sigma_X| |I + \Sigma_X^{-1} H_X \Sigma_M H_X^T|}{|\Sigma_X|} \right) \quad (48)$$

$$= \frac{1}{2} \log |I + \Sigma_X^{-1} H_X \Sigma_M H_X^T|. \quad (49)$$

Now, suppose for the sake of contradiction that $H_X^T \Sigma_X^{-1} H_X \not\succcurlyeq H_Y^T \Sigma_Y^{-1} H_Y$. Then, by the definition of positive semidefiniteness, $\exists c \in \mathbb{R}^{d_M}$ such that

$$c^T H_X^T \Sigma_X^{-1} H_X c < c^T H_Y^T \Sigma_Y^{-1} H_Y c. \quad (50)$$

Since $\log|I + AB| = \log|I + BA|$, $\log(\cdot)$ is an increasing function, and the determinant of a scalar is equal to itself, we have that

$$\log|1 + c^T H_X^T \Sigma_X^{-1} H_X c| < \log|1 + c^T H_Y^T \Sigma_Y^{-1} H_Y c| \quad (51)$$

$$\Rightarrow \log|I + \Sigma_X^{-1} H_X c c^T H_X^T| < \log|I + \Sigma_Y^{-1} H_Y c c^T H_Y^T| \quad (52)$$

Now, since $cc^T \succcurlyeq 0$, it is a valid covariance matrix. So if we set $\Sigma_M \triangleq cc^T$, we get

$$I(M; X) = \log|I + \Sigma_X^{-1} H_X \Sigma_M H_X^T| \quad (53)$$

$$< \log|I + \Sigma_Y^{-1} H_Y \Sigma_M H_Y^T| \quad (54)$$

$$= I(M; Y). \quad (55)$$

However, this contradicts (44), which holds no matter what Σ_M is. Therefore, we must have that $H_X^T \Sigma_X^{-1} H_X \succcurlyeq H_Y^T \Sigma_Y^{-1} H_Y$. \square

APPENDIX C

THE APPROXIMATE DEFICIENCY IS WELL DEFINED

Lemma 7. For jointly Gaussian random vectors M , X and Y satisfying the assumption in Remark 1, \hat{T} and $\hat{\Sigma}_T$ as defined by (17) and (18) yield an approximate deficiency $\hat{\delta}_G(M : Y \setminus X)$ that is well defined, i.e. the covariance matrix $\hat{\Sigma}_T + \hat{T}\hat{T}^T$ for the composite channel $Q_{Y|X} \circ P_{X|M}$ is invertible.

Proof. We know that $\hat{\Sigma}_T \succcurlyeq 0$ by virtue of the constraint in (18) and thus $\hat{\Sigma}_T + \hat{T}\hat{T}^T \succcurlyeq 0$ as well. For ease of notation, let $A \triangleq H_Y \Sigma_M H_Y^T$ and $B \triangleq H_X \Sigma_M H_X^T$. Assume for a

contradiction that $\hat{\Sigma}_T + \hat{T}\hat{T}^T$ is rank deficient. This implies that there exists a vector $v \neq 0$ such that:

$$v^T (\hat{\Sigma}_T + \hat{T}\hat{T}^T) v = v^T (I + A - \hat{T}B\hat{T}^T) v = 0$$

$$\Rightarrow v^T (I + A) v = v^T (\hat{T}B\hat{T}^T) v$$

$$\stackrel{(a)}{\Rightarrow} (\hat{T}^T v)^T B (\hat{T}^T v) > 0$$

$$\Rightarrow \hat{T}^T v \neq 0$$

where (a) follows from $A = (H_Y \Sigma_M^{-\frac{1}{2}})(H_Y \Sigma_M^{-\frac{1}{2}})^T \Rightarrow A \succcurlyeq 0$. But since $\hat{\Sigma}_T \succcurlyeq 0$, we have:

$$v^T \hat{\Sigma}_T v \geq 0$$

$$\Rightarrow v^T (\hat{\Sigma}_T + \hat{T}\hat{T}^T) v - v^T (\hat{T}\hat{T}^T) v \geq 0$$

$$\Rightarrow v^T \hat{T}\hat{T}^T v = \|\hat{T}^T v\|_2^2 \leq 0.$$

which is a contradiction since $\hat{T}^T v \neq 0 \Rightarrow \|\hat{T}^T v\|_2^2 > 0$. \square

APPENDIX D

PROOF OF PROPOSITION 4

For ease of notation, we omit $(M : Y \setminus X)$ and refer simply to $\hat{\delta}_G$, δ_G , and δ .

($\delta = 0 \Rightarrow \hat{\delta}_G = 0$) It was shown by Torgersen [13, Theorem 8.2.13] that $\delta = 0 \Rightarrow \delta_G = 0$, thus it suffices to show that $\delta_G = 0 \Rightarrow \hat{\delta}_G = 0$. Since $\delta_G = 0$ implies that there exists a T^* such that $T^* H_X = H_Y$, we know that the objective function in (18) is minimized by $T = T^*$. To show that T^* is in the feasible set, we note that $\delta_G = 0$ implies that there exists a $\Sigma_T^* \succcurlyeq 0$ such that $\Sigma_T^* + T^* T^{*T} = I \Rightarrow I - T^* T^{*T} \succcurlyeq 0$. As such, when $T = T^*$ we have that $\Sigma_T = I + H_Y \Sigma_M H_Y^T - T^* (I + H_X \Sigma_M H_X^T) T^{*T} = I - T^* T^{*T} = \Sigma_T^*$.

($\delta = 0 \Leftarrow \hat{\delta}_G = 0$) This follows from $\hat{\delta}_G \geq \delta \geq 0$.

($\delta = I(M; Y) \Rightarrow \hat{\delta}_G = I(M; Y)$) Note that if $\hat{T} = 0$, we have $\hat{\Sigma}_T = I + H_Y \Sigma_M H_Y^T \Rightarrow \tilde{Q}_{Y|X} \circ P_{X|M} = P_Y \Rightarrow \hat{\delta}_G = I(M; Y)$. As such, it suffices to show that $\delta = I(M; Y)$ implies that the objective function in (18) is minimized at $T = 0$. Suppose for a contradiction that there exists a \tilde{T} such that:

$$\mathbb{E}_{P_M} [\|(\tilde{T} H_X - H_Y) M\|_A^2] < \mathbb{E}_{P_M} [\|H_Y M\|_A^2] \quad (56)$$

$$\Rightarrow \mathbb{E}_{P_M} [\|(\lambda \tilde{T} H_X - H_Y) M\|_A^2] < \mathbb{E}_{P_M} [\|H_Y M\|_A^2]$$

for any $\lambda \in (0, 1)$ and with $A \triangleq I + H_Y \Sigma_M H_Y^T$ giving the marginal covariance matrix for Y (i.e. $P_Y = \mathcal{N}(0, A)$). The implication follows from the convexity of the Mahalanobis distance, noting that $\lambda \tilde{T} = \lambda \tilde{T} + (1 - \lambda)0$. As such, we can assume without loss of generality that \tilde{T} can be chosen such that (56) holds and $\tilde{\Sigma}_T \triangleq A - \tilde{T}\tilde{T}^T$ is PSD. Defining $\tilde{Q}_{Y|X} \triangleq \mathcal{N}(\tilde{T} X, \tilde{\Sigma}_T)$, we note that $\tilde{Q}_{Y|X} \circ P_{X|M} = \mathcal{N}(\tilde{T} H_X M, A)$. As such:

$$\mathbb{E}_{P_M} [D(P_{Y|M} \| \tilde{Q}_{Y|X} \circ P_{X|M})] - I(M; Y)$$

$$= \mathbb{E}_{P_M} [D(P_{Y|M} \| \tilde{Q}_{Y|X} \circ P_{X|M}) - D(P_{Y|M} \| P_Y)]$$

$$= \mathbb{E}_{P_M} [\|(\tilde{T} H_X - H_Y) M\|_A^2] - \mathbb{E}_{P_M} [\|H_Y M\|_A^2] < 0.$$

But since $\tilde{Q}_{Y|X} \in \mathcal{C}_G(Y | X)$, we have:

$$\delta_G \leq \mathbb{E}_{P_M} \left[D(P_{Y|M} \| \tilde{Q}_{Y|X} \circ P_{X|M}) \right] < I(M; Y) \leq \delta, \quad (57)$$

which is a contradiction. \square

APPENDIX E

DETAILS ON CVX IMPLEMENTATION

We briefly discuss how to reformulate the optimization problem in (18) such that it satisfies the disciplined convex programming (DCP) rules and can be solved using the CVX software package. First we note that the objective function can be rewritten using the trace trick. Letting $A = I + H_Y \Sigma_M H_Y^T$ for ease of notation:

$$\begin{aligned} & \mathbb{E}_{P_M} [\| (TH_X - H_Y)M \|_A^2] \\ &= \mathbb{E}_{P_M} [M^T (TH_X - H_Y)^T A^{-1} (TH_X - H_Y) M] \\ &= \mathbb{E}_{P_M} [\text{Tr} (M^T (TH_X - H_Y)^T A^{-1} (TH_X - H_Y) M)] \\ &= \mathbb{E}_{P_M} [\text{Tr} (M M^T (TH_X - H_Y)^T A^{-1} (TH_X - H_Y))] \\ &= \text{Tr} (\Sigma_M (TH_X - H_Y)^T A^{-1} (TH_X - H_Y)) \\ &= \| A^{-\frac{1}{2}} TH_X \Sigma_M^{\frac{1}{2}} - A^{-\frac{1}{2}} H_Y \Sigma_M^{\frac{1}{2}} \|_F^2 \end{aligned}$$

where $\|\cdot\|_F^2$ is the squared Frobenius norm. Next, the constraint can be rewritten using the Schur complement [17]:

$$\begin{aligned} I + H_Y \Sigma_M H_Y^T - T(I + H_X \Sigma_M H_X^T) T^T &\succcurlyeq 0 \\ \Updownarrow \\ \begin{bmatrix} I + H_Y \Sigma_M H_Y^T & T \\ T^T & (I + H_X \Sigma_M H_X^T)^{-1} \end{bmatrix} &\succcurlyeq 0 \end{aligned}$$

Combining the two, we obtain the problem in a form that can be solved directly by CVX:

$$\begin{aligned} \hat{T} = \underset{T}{\text{argmin}} \quad & \| A^{-\frac{1}{2}} TH_X \Sigma_M^{\frac{1}{2}} - A^{-\frac{1}{2}} H_Y \Sigma_M^{\frac{1}{2}} \|_F^2 \\ \text{s.t.} \quad & \begin{bmatrix} I + H_Y \Sigma_M H_Y^T & T \\ T^T & (I + H_X \Sigma_M H_X^T)^{-1} \end{bmatrix} \succcurlyeq 0 \end{aligned} \quad (58)$$

APPENDIX F

STEP-BY-STEP EXPERIMENTAL PROCEDURE

- 1) Sample d_M, d_X, d_Y according one of (S1)–(S4).
- 2) Sample a covariance matrix $\Sigma \in \mathbb{R}^{d \times d}$ (where $d = d_M + d_X + d_Y$) from a standard Wishart distribution.
- 3) Compute $I(M; (X, Y))$, $I(M; X)$, and $I(M; Y)$.
- 4) Compute conditional mean and covariances and whiten to obtain H_X, H_Y , and $\Sigma_X = \Sigma_Y = I$.
- 5) Estimate the deficiencies $\hat{\delta}_G(M : Y \setminus X)$ and $\hat{\delta}_G(M : X \setminus Y)$, using the the Python CVX package to solve (58).
- 6) Compute the approximate δ -PID atoms.
- 7) Check if \widehat{RI} and \widehat{SI} are non-negative.
- 8) If $d_M = 1$, check if either $\widehat{UI}_X \approx 0$ or $\widehat{UI}_Y \approx 0$.

APPENDIX G
3-DIMENSIONAL SIMPLEX VIEWS




Fig. 2: Alternative views of Figure 1. Each row displays one of the four sampling schemes, and each column provides a rotated view of the simplex.